

OVIA Incentive Demo

What You Need

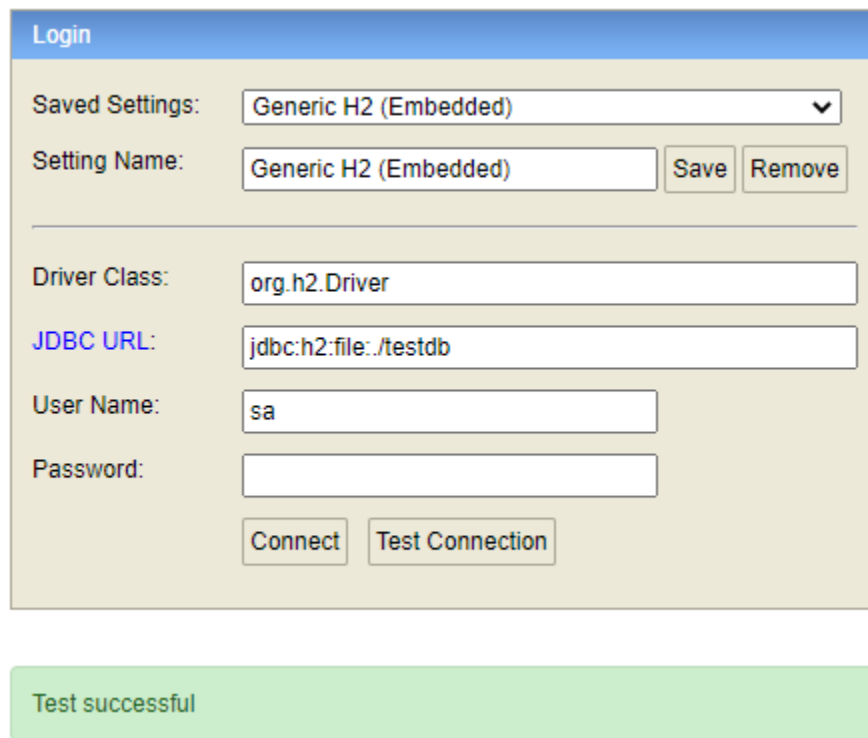
- JDK 1.8 or later
- Maven 3.2+
- Favorite IDE

Download

- git clone <https://github.com/t2r/ovia.git>

How to run the demo

- cd ovia-incentive-demo
- mvn clean install
- mvn spring-boot:run (server)
- ./client-test.sh (client requests to trigger awards)
 - This is just using *curl* to spend get/post requests to the REST API endpoints.
- You may also query the database tables directly by hitting <http://localhost:9000/h2>

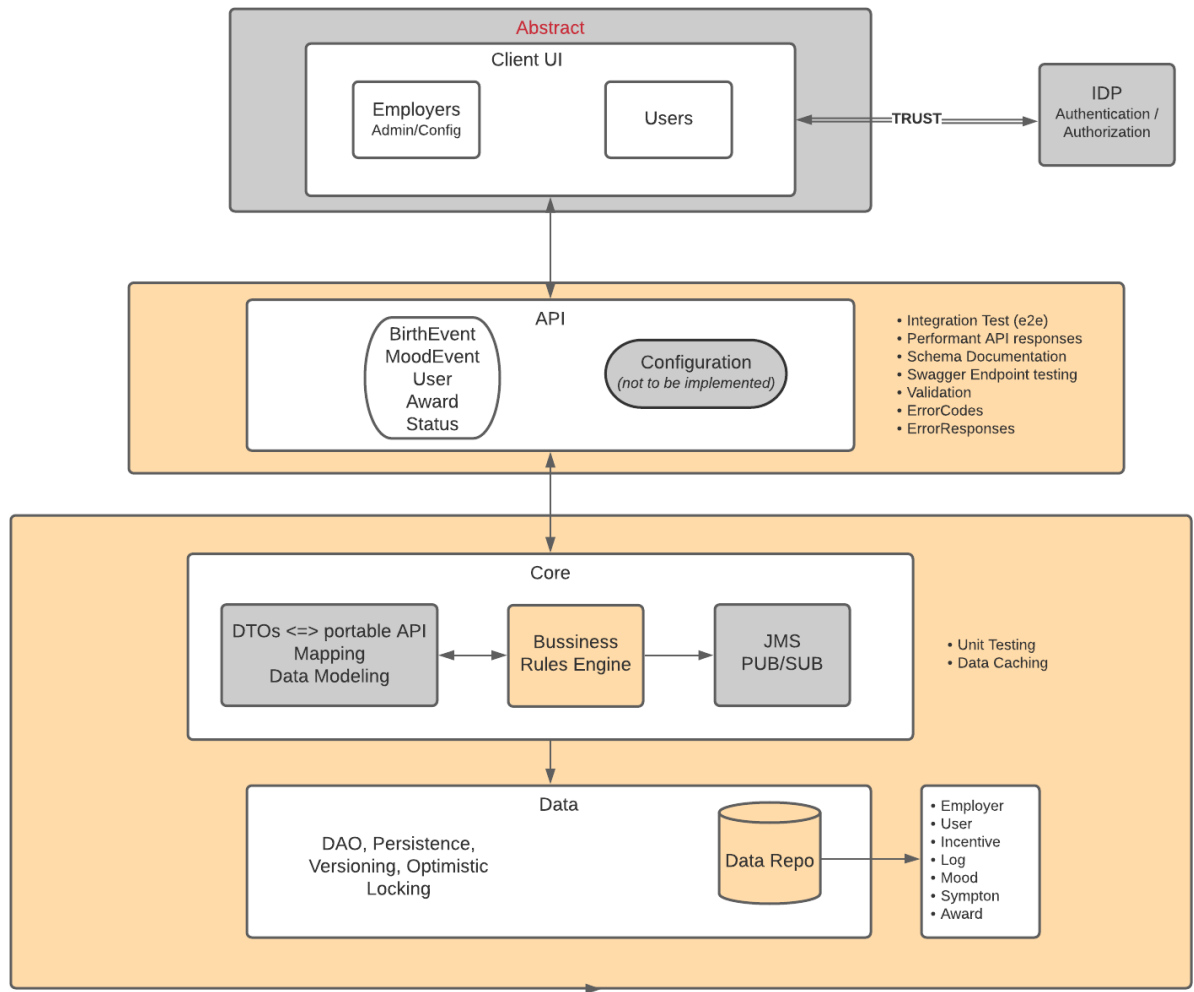


The screenshot displays the H2 database console interface. At the top, there is a 'Login' section with a dropdown menu for 'Saved Settings' set to 'Generic H2 (Embedded)'. Below this, the 'Setting Name' is also 'Generic H2 (Embedded)', with 'Save' and 'Remove' buttons. The main configuration area includes fields for 'Driver Class' (org.h2.Driver), 'JDBC URL' (jdbc:h2:file:./testdb), 'User Name' (sa), and 'Password'. At the bottom of this section are 'Connect' and 'Test Connection' buttons. A green banner at the very bottom of the image indicates 'Test successful'.

Design Overview

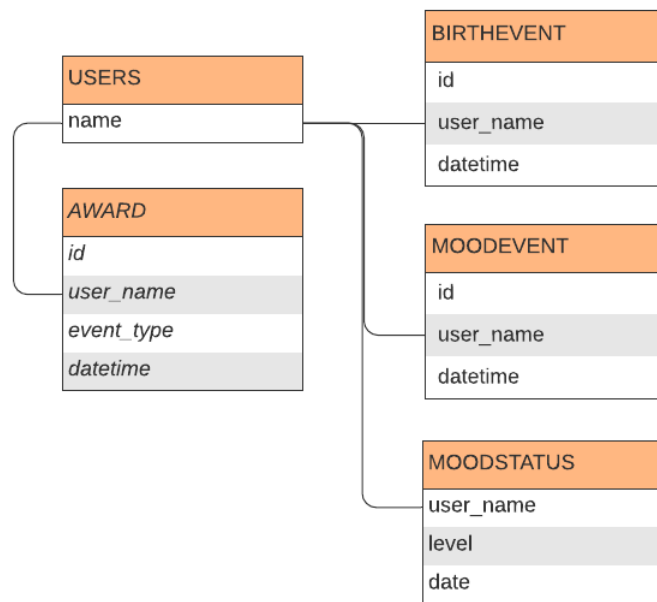
- I first started with an overview software component architecture design to determine how I was going to break this task down into a manageable demo and to help identify the components/frameworks I would need to accomplish this task as well as identify areas that that

would not be able to focus on.



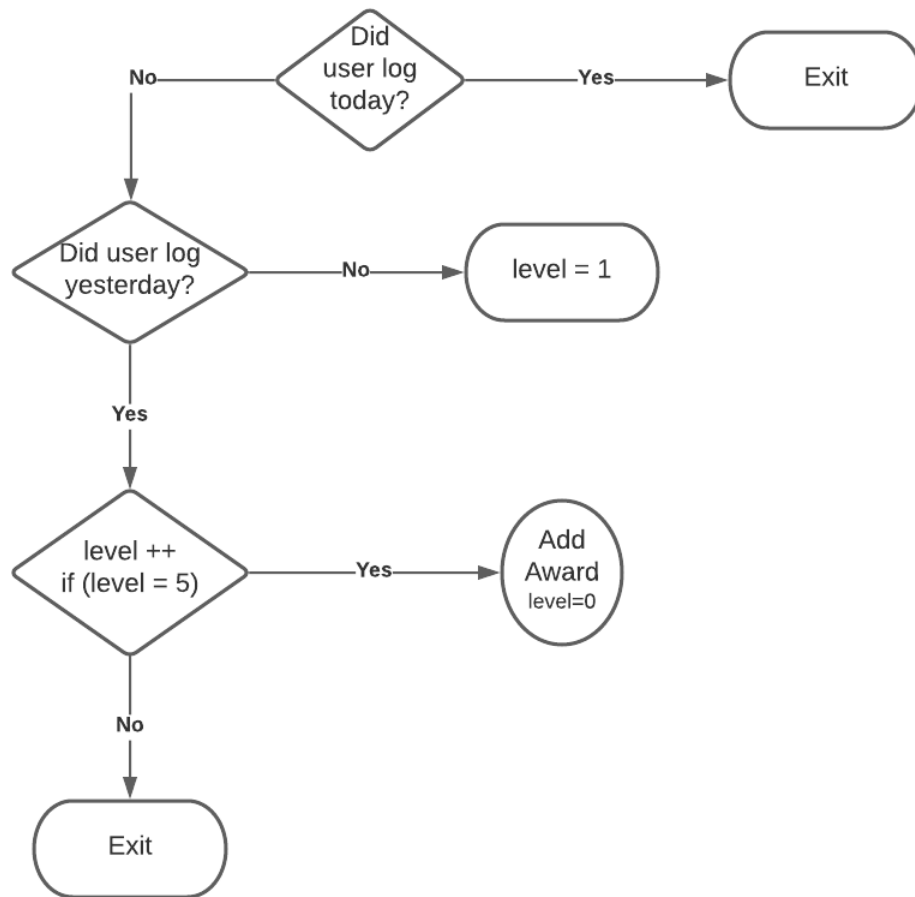
1. I knew I would not get to get documentation, testing, mapping from DTO to portable JSON API objects, flexible dynamic configuration, a *real* BRE to trigger events, etc. Of course, BRE, is one of the most significant performance improvements that could be considered here. The demo is simplistic in that it is single-threaded and executes all logic before returning response back to client. My thoughts immediately shifted to concurrency, data storage size where perhaps a background scanner that is scheduled to run off-peak hours/etc to collect statistics and determine awards without blocking the API response/client and send push notifications instead would be far much better approach.

- My next step, since I had no real system integration requirements, was to determine the data layer needed to help conceptualize the api layer that would be needed.



- I then knew I had two algorithm requirements to meet:
 1. Reported BIRTH event, which is just a direct relationship with earning an award.

2. User logged data five days in a row. For this, I created a state transition diagram to help me decide how I would accomplish this code wise:



- I knew that that algorithm implementation would be extremely rigid meaning that I wasn't going to be able to dive into a flexible, dynamic configuration based on customer configuration/requirements that would drive the behavior for how this functions.

Improvements

- Documentation - code, API schema
- Unit Testing
 - Something like swagger to allow for interactive API testing and viewing documentation.
 - End-to-End integration tests that test the entire work-flow.
 - Additional unit tests that provide more code coverage with min, max, edge-case, and bad data scenarios.
- A BRE system that allows event triggering/scheduling to decouple heavy algo processing from user API response time.
 - Something like a PUB/SUB model came to mind where notifications could be pushed.
- Concurrency, optimistic locking, versioning

- Configuration, business rules, business policy were surrounding impacts that did not contain enough requirements to properly address.

Data Sharing

- This needs a lot more surrounding requirements to give it proper attention. But here are my initial thoughts:
 - IDP and trust
 - Reporting Subsystem. Rather than data access, are the customer's simply look for metrics and data reporting charts/views that can viewed and/or exported as PDF's, for example.